



Kissipo Learning for Deep Learning

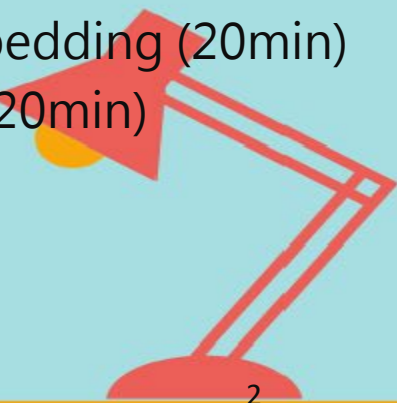
Topic 18: AOI simple Pipeline (A) (20min)

Hsueh-Ting Chu

KLDL-W7-T18

Topics

- Topic 01: Introduction to Deep Learning (20min)
- Topic 02: Kissipo Learning for Deep Learning (20min)
- Topic 03: Python quick tutorial (20min)
- Topic 04: Numpy quick tutorial (15min)
- Topic 05: Pandas quick tutorial (15min)
- Topic 06: Scikit-learn quick tutorial (15min)
- Topic 07: OpenCV quick tutorial (15min)
- Topic 08: Image Processing basics (20min)
- Topic 09: Machine Learning basics (20min)
- Topic 10: Deep Learning basics (20min)
- Topic 11: TensorFlow overview (20min)
- Topic 12: CNN with TensorFlow (20min)
- Topic 13: RNN with TensorFlow (20min)
- Topic 14: PyTorch overview (20min)
- Topic 15: CNN with PyTorch (20min)
- Topic 16: RNN with Pytorch (20min)
- Topic 17: Introduction to AOI (20min)
- **Topic 18: AOI simple Pipeline (A) (20min)**
- Topic 19: AOI simple Pipeline (B) (20min)
- Topic 20: Introduction to Object detection (20min)
- Topic 21: YoloV5 Quick Tutorial (20min)
- Topic 22: Using YoloV5 for RSD (20min)
- Topic 23: Introduction to NLP (20min)
- Topic 24: Introduction to Word Embedding (20min)
- Topic 25: Name prediction project (20min)



Course Schedule

- W1 - Course Introduction
- W2 - DL Programming Basics(1)
- W3 - DL Programming Basics(2)
- W4 - DL with TensorFlow
- W5 - Midterm
- W6 - DL with PyTorch
- W7 - AOI hands-on project
- W8 - RSD hands-on project
- W9 - NLP hands-on project
- W10 - Final exam

DL: Deep Learning

AOI: Automated Optical Inspection

RSD: Road Sign Detection

NLP: Natural Language Processing





Week 7 Topics

- Topic 17: Introduction to AOI (20min)
- Topic 18: AOI simple Pipeline (A) (20min)
- Topic 19: AOI simple Pipeline (B) (20min)



AOI pipeline (A)



 KLDL-18-AOI simple Pipeline (A) ☆

檔案 編輯 檢視畫面 插入 執行階段 工具 說明 已儲存所有變更

目錄

Topic 18: AOI simple Pipeline (A)

Exercise: Full solution

Step 1: Load the dataset from google drive

Step 2: Import PyTorch libraries

Step 3: read the training set

Step 4: Show statistics of training images

Step 5: Choose one of CNN models

Step 6: Instancing a dataloader

Step 7: Set up a train dataloader with a custom dataset

Step 8: total_batch


Step 9: Train model

Step 10: Save the trained model

Step 11: Check training results

Step 12: Analyze training results

+ 程式碼 + 文字



Deep Learning Course

Topic 18: AOI simple Pipeline (A)

▼ Exercise: Full solution

- Single CNN model
- ImageDataSet
- ImageDataLoader
- Submit results

Aldea AOI Project <https://aidea-web.tw/topic/285ef3be-44eb-43dd-85cc-f0388bf85ea4>

▼ Step 1: Load the dataset from google drive

Step 1: Load the dataset from google drive



AUAOI Ex2.

CO PRO KLDL-18-AOI simple Pipeline (A) ☆

檔案 編輯 檢視畫面 插入 執行階段 工具 說明 已儲存所有變更

+ 程式碼 + 文字

Exercise: Full solution

- Single CNN model
- ImageDataSet
- ImageDataLoader
- Submit results


Aldea AOI Project <https://aidea-web.tw/topic/285ef3be-44eb-43dd-85cc-f0388bf85ea4>

Step 1: Load the dataset from google drive

If the following command does not work, please download it, put it on your Google drive, and set up sharing

Download from: <https://drive.google.com/file/d/1tovCO2gsjesjJ80sfHgahyt-buY34dk0/view?usp=sharing>

```
[ ] %%bash
pip install --upgrade --no-cache-dir gdown
gdown https://drive.google.com/uc?id=1tovCO2gsjesjJ80sfHgahyt-buY34dk0
unzip aoi-dataset.zip
rm aoi-dataset.zip
```




CO aidea_aoi_ex1.ipynb ☆

File Edit View Insert Runtime

Files

- sample_data
- test_images
- train_images
- content
- test.csv
- train.csv

Step 2: Import python libraries



AUAOI Ex2.

```
[ ] import os
import glob
import torch
from torch import nn
from torch.utils.data import Dataset, DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor
```


```
[ ] print (torch.cuda.is_available())
```

```
[ ] device_name=torch.cuda.get_device_name(0)
print(f"Using GPU {device_name}")
```

```
[ ] import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```



Step 3: read the training set



AUAOI Ex2.

```
[ ] import pandas as pd
    df_train = pd.read_csv("train.csv")
    print(df_train.shape)

[ ] df_train.head()

[ ] train_files = df_train.iloc[:,0].values
    train_labels = df_train.iloc[:,1].values
    print(train_labels[:10])
```




```
[6] df_train.head()
```

	ID	Label
0	train_00000.png	0
1	train_00001.png	1
2	train_00002.png	1
3	train_00003.png	5
4	train_00004.png	5

```
[7] train_files = df_train.iloc[:,0].values
    train_labels = df_train.iloc[:,1].values
    print(train_labels[:10])
```

```
[ ] ['0' '1' '1' '5' '5' '5' '3' '0' '3' '5']
```



Step 4: Show statistics of training images



AUAOI Ex2.

```
import seaborn as sns  
g = sns.countplot(train_labels)
```

```
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py  
import pandas.util.testing as tm
```



Step 5: Choose one of CNN models



AUAOI Ex2.

EfficientNet B0 to B7

Model-EfficientNet

https://pytorch.org/hub/nvidia_deeplearningexamples_efficientnet/

Base model	resolution	Base model	resolution
EfficientNetB0	224	EfficientNetB4	380
EfficientNetB1	240	EfficientNetB5	456
EfficientNetB2	260	EfficientNetB6	528
EfficientNetB3	300	EfficientNetB7	600

```
[ ] num_classes=6
```

```
[ ] modelfile = None
```

```
[ ] import torchvision.models as models
    model=models.efficientnet_b0(num_classes=num_classes)
    if modelfile != None: model.load_state_dict(torch.load(modelfile))
    model.cuda()
```



Step 6: Instancing a dataloader



AUAOI Ex2.

- Transforms
- CustomDataset
- dataloader

```
from torchvision import transforms
pretrained_size = 224
pretrained_means = [0.485, 0.456, 0.406]
pretrained_stds = [0.229, 0.224, 0.225]
train_transform = transforms.Compose([
    transforms.Resize(pretrained_size),
    transforms.RandomRotation(5),
    transforms.RandomHorizontalFlip(0.5),
    transforms.RandomCrop(pretrained_size, padding=4),
    transforms.ToTensor(),
    transforms.Normalize(mean=pretrained_means, std=pretrained_stds)
])
```

```
from PIL import Image
class CustomDataset(torch.utils.data.Dataset):
    def __init__(self, csv_path, images_folder, transform = None):
        self.df = pd.read_csv(csv_path)
        self.images_folder = images_folder
        self.transform = transform

    def __len__(self):
        return len(self.df)

    def __getitem__(self, index):
        filename = self.df.iloc[index]['ID']
        label = self.df.iloc[index]['Label']
        image = Image.open(os.path.join(self.images_folder, filename))
        if self.transform is not None:
            image = self.transform(image)
        return image, label
```

Step 7: Set up a train dataloader with a custom dataset




AUAOI Ex2.

```
▶ batches = 24
imgdir= "train_images"
csvfile = "train.csv"

[ ] train_dataset = CustomDataset(csvfile, imgdir, train_transform)
train_dataloader = DataLoader(train_dataset, batch_size=batches, shuffle=True)
print(f"Total images={len(train_dataset)}")
```



Step 8: total_batch



AUAOI Ex2.

```
total_batch=len(train_dataset)//batches + 1  
print(total_batch)
```



Step 9: Train model



AUAOI Ex2.

```
▶ loss = nn.CrossEntropyLoss()
  optimizer = torch.optim.SGD(model.parameters(), lr=0.001, momentum=0.9)


[ ] num_epochs = 10

[ ] for epoch in range(num_epochs):
    for i, (batch_images, batch_labels) in enumerate(train_dataloader):
        # Zero your gradients for every batch!
        optimizer.zero_grad()
        inputs = batch_images.cuda()
        labels = batch_labels.cuda()
        # Make predictions for this batch
        outputs = model(inputs)
        # Compute the loss and its gradients
        cost = loss(outputs, labels)
        cost.backward()
        # Adjust learning weights
        optimizer.step()

    if (i+1) % 100 == 0:
        print(f'Epoch [{epoch+1}/{num_epochs}], lter [{i+1}/{total_batch}]')
```



Step 10: Save the trained model




AUAOI Ex2.

```
filepath = "AOI-EnBO.pth"  
torch.save(model.state_dict(), filepath)
```



Step 11: Check training results




AUAOI Ex2.

```
model.eval()
# again no gradients needed
with torch.no_grad():
    total_batch = len(test_dataset)//batches
    for i, (batch_images, batch_labels) in enumerate(test_dataloader):
        images = batch_images.cuda()
        labels = batch_labels.cuda()
        outputs = model(images)
        _, predictions = torch.max(outputs, 1)
        train_predictions[i*batches:(i+1)*batches] = predictions.cpu()
        if (i+1) % 10 == 0:
            print(f'lter [{i+1}/{total_batch}]')

train_predictions=train_predictions.astype(int)
```



Step 12: Analyze training results



AUAOI Ex2.

```
from sklearn.metrics import confusion_matrix
confusion=confusion_matrix(train_labels, train_predictions)
print(confusion)

train_num = 2528
overkill= []
underkill = []
for i,(label, prediction) in enumerate(zip(train_labels, train_predictions)):
    if label == 0 and prediction !=0:
        overkill.append(i)
    if label != 0 and prediction ==0:
        underkill.append(i)
print('# of overkill= {}; # of underkill= {}'.format(len(overkill), len(underkill)))
```





Thanks!

Q&A

